

The Choice is Yours

Below are some links to some of the API Workbenches mentioned

- Anthropic - console.anthropic.com
- OpenAI - platform.openai.com
- Google - ai.google.dev
- Mistral - console.mistral.ai

Anatomy of an AI Request

How building in the Workbench looks

- Prompt Library System: Most Workbenches will allow you to save your prompts to your account for easy testing. Figure out a naming system early on to account for high volume of prompts.
- Model Selection: Click 'claude-sonnet-4-5...' to switch Anthropic models.
- System Prompt: This is the static instruction you're sending to the LLM, this defines what this prompt fundamentally does.
- User Request: 'User' can be misleading here for non-chatbot systems. This is just the particular data you're sending for this particular instance.
- Expected Response: "This is an LLM response from Anthropic API"

The screenshot shows the Cohere Workbench interface. At the top, there are three buttons: a list icon, a plus icon, and a 'Default' dropdown. Below that, a header bar displays 'Testing New System' with a lock icon and the text 'Last saved Nov 16 at 10:59 AM'. To the right of the header are buttons for 'claude-sonnet-4-5-20250929', a file icon, 'Examples', and 'Templatize'. The main area is divided into two sections: 'System Prompt' and 'User'. The 'System Prompt' section contains the text: 'This is an example prompt, just print out "This is an LLM response from Anthropic API" no matter what I send to you.' The 'User' section contains the text: 'Trying this out!'. At the bottom of the interface are buttons for 'Pre-fill response' and 'Add message pair'. The bottom right corner of the slide has a page number '43'.