

Quick Reminder...

For people who may have forgotten

- This course isn't for coding. However, there will be a code repo linked in the description that has common code patterns that you will likely need for calling third-party LLM APIs, parsing responses, counting tokens and assessing cost, etc.
- If you want a more focused learning environment, take the course on nicholasmgoldstein.com/courses. This way you won't be tempted to jump to a recommended video or divert your attention
- This is a free course, but if you want to support me, you can like/subscribe or check out some of my own AI platforms like Emstrata @ emstrata.com
- I'm fully open to your input. If there's topics you'd like covered more thoroughly or anything you're unclear on, feel free to comment

Picking an LLM

Making your choice early on prevents paralysis later

- Different LLMs will have different strengths and textual voice. You will learn this best by working with them and iterating on system prompts, rather than studying benchmarks
- My personal favorite is Claude/Anthropic's third-party API (3.7 => 4.0), but there are other options like ChatGPT/OpenAI or Gemini/Google. You will need to add a credit card to use their platforms, as they will charge for usage. Pricing is found on their websites as well
- You may opt for a couple LLMs for different types of tasks. For example, Claude 4.0 for heavier analysis and rule-following, but Mistral or another smaller model for simpler tasks. This would allow you to save money and retain most of your effectiveness
- A benefit to using OpenAI or Google over Anthropic is they have other tools built into their platforms, like TTS, transcription, image generation, etc.