Reasoning/Strategy Layers

The decision-maker of your platform

- Reasoning layers make decisions before content gets generated. They evaluate the current state, consider available options, assess consequences, and choose a direction. Think of them as the "planning brain" of your system

- In Emstrata, Discovery handles this - it looks at what the participant wants to do, considers the simulation state, evaluates what outcomes make narrative sense, and determines how the action should resolve. It's not writing the story yet; it's deciding what should happen.

- **When you need one:** If you find yourself asking an LLM to both "figure out what should happen AND write it beautifully," you're overloading a single prompt. Split it. Reason first, write second.

Memory Consolidation Layers

The stenographer of your platform

- Memory consolidation layers distill what just happened into something retrievable later. They extract the important details from verbose content and store them in a format your system can efficiently query or format into future inputs

- In Emstrata, Groundskeeper serves this function. After Discovery determines what happens and Narration writes it, Groundskeeper updates the comprehensive memory of all entities and the emergent narrative. It’s maintaining the source of truth about the simulation state