Content Layers

The performer of your platform

- Content layers generate the actual output users experience - the prose, dialogue, descriptions, or interface text. These layers take decisions from reasoning layers and context from memory layers, then craft the experience

- Emstrata's Narration layer does this. It receives Discovery's decisions about what happened, checks Groundskeeper's simulation state, and writes the actual narrative text that players read. It's optimizing for atmosphere, pacing, and emotional resonance - not logic or consistency (that's handled elsewhere)

Catch-All/Connector Layers

The clean-up crew of your platform

- Not every layer fits a clean category. Catch-all layers are hybrids that do complementary work for multiple other layers. They handle tasks that don't belong to any single specialized layer but are essential for the system to function cohesively

- These layers often emerge when you discover gaps like two layers need to work together but speak different "languages," or several layers all need the same preprocessing that none of them should be responsible for individually

- In Emstrata, the Chron-Con does more than just error correction. It also tracks secrets and memories from the narrative, explicitly tagging them for Groundskeeper to integrate into system memory. You don't want Narration burdened with the unrelated task of extracting and categorizing secrets while it's trying to write high-quality prose. And Groundskeeper needs these pieces explicitly labeled as "secrets" or "memories" to properly integrate them into the simulation history. The Chron-Con bridges this gap.